MIGRATION STATUS REPORT

Airtable to Salesforce: what loaded, what didn't, and why

A reconciliation of every Airtable table against the records now in the Salesforce sandbox, the transformations applied along the way, and the decisions still needed from the Partnerships team.

Environment gsa-peo sandbox **Report date** 13 August 2026 **Source** Airtable, pulled 12 August 2026

Where things stand

Ten of the eleven data sets now hold records in Salesforce; Meetings has not been started. Every record that could be loaded without guessing at ambiguous data has been. The gaps below are almost entirely traceable to a handful of source-data issues, most of which release a large number of records at once.

6,735	10 / 11	169	1,845
Records created or updated in Salesforce	Data sets loaded, in whole or in part	Unmatched accounts — the largest single blocker	Meetings not yet started

THE ONE NUMBER TO WATCH

A single issue — **169 Airtable account rows that don't match an account already in Salesforce** — is responsible for the majority of every other gap on this page. It blocks partner accounts, which blocks applications, which blocks the application–contact links. Resolving those 169 rows would pull several hundred additional records into Salesforce with no further engineering work.

Record reconciliation

Airtable row counts against what is now in Salesforce. A negative difference is not necessarily data loss — in several cases it is intentional consolidation, explained in the final column.

DATA SET	AIRTABLE	SALESFORCE	DIFF	%	STATUS	EXPLANATION
Market Segments	6	6	0	100%	COMPLETE	Loaded previously; untouched by this migration.

DATA SET	AIRTABLE	SALESFORCE	DIFF	%	STATUS	EXPLANATION
Accounts	757	588	−169	77.7%	NEEDS DECISIONS	Accounts already existed in Salesforce, so these were matched rather than created. 169 Airtable rows matched nothing. Spot-checking found most are duplicate rows *within Airtable* — e.g. “Army” alongside “Department of the Army”.
Partner Accounts	99	74	−25	74.7%	NEEDS DECISIONS	20 are blocked by the unmatched accounts above. 5 have no usable account link at all (4 marked inactive, 1 linked to two accounts).
Contacts	2,119	1,936	−183	91.4%	LOADED	Source is two tables: 1,599 Contacts + 520 Opportunity Contacts. The difference is mostly *intentional* — the same person appearing on several rows was consolidated into one contact. 45 rows had neither a name nor an email and could not be loaded; 8 were rejected as duplicates of existing Salesforce contacts.
Opportunities	928	742	−186	80.0%	NEEDS DECISIONS	142 blocked by unmatched accounts, 28 have no status recorded, 16 have no account link.

DATA SET	AIRTABLE	SALESFORCE	DIFF	%	STATUS	EXPLANATION
Applications	1,064	688	−376	64.7%	NEEDS DECISIONS	359 are blocked because their partner account hasn't loaded — again tracing back to the unmatched accounts. 17 have no partner account link (6 retired applications, 11 in-progress drafts).
Impediments	41	39	−2	95.1%	LOADED	2 rows are entirely blank and look like accidental empty records.
Application ↔ Contact links	2,797	1,880	−917	67.2%	PARTIAL	A link can only exist once both sides do; 849 are waiting on an application that hasn't loaded. 33 were duplicates created by consolidating contacts and were merged.
Impediment ↔ Opportunity links	783	267	−516	34.1%	NEEDS DECISIONS	465 of these point at a single impediment named “None”, which we believe is a placeholder meaning *no impediment* — see below. 44 more await an opportunity that hasn't loaded.
Opportunity contact roles	520	515	−5	99.0%	LOADED	437 of 520 source rows loaded, producing 515 role records because some people hold more than one role. 83 rows await an opportunity or contact.
Meetings	1,845	0	−1,845	0%	NOT STARTED	Deferred — Airtable holds no meeting times. See “Known issue: meetings cannot be migrated as calendar events”.

The contacts finding

This is the most consequential thing we learned about the source data, and it changes how contacts should be maintained in Airtable going forward.

The contacts attached to opportunities are a separate population

Airtable has a **Contacts** table (1,599 rows) and, separately, an **Opportunity Contacts** table (520 rows). We expected the second to reference the first. It does not — it has no link to the Contacts table at all, only a typed name and an email address.

When we checked those 520 people against the Contacts table:

- **65** existed in the Contacts table.
- **348** did not appear there at all.
- **107** had no email address recorded, only a name.

In other words, roughly two-thirds of the people attached to opportunities existed nowhere else in the base. They had to be created as contacts in their own right before their roles on opportunities could be recorded — a separate load, not a lookup.

> WHY THIS MATTERS BEYOND THE MIGRATION
>
> Salesforce treats a person as one record who can be attached to many opportunities and applications. Airtable currently has no equivalent link, so the same person gets typed in again for each thing they're involved with. That's why the Contacts table also contains duplicates: **61 email addresses appear on two or more rows**, one of them four times. We consolidated those into single contacts on load, but the underlying pattern will keep producing duplicates until Airtable has a proper way to link one person to several applications.

A related case: the same person, several roles

62 of the opportunity contact rows list more than one role for the same person on the same opportunity — for example *Senior POC + Decision Maker + Day-to-Day POC*. Salesforce stores one role per record, so this needed an explicit decision rather than a default.

Decision: create one record per role

The alternative was to pick a single “most senior” role and discard the others, which would have silently lost information the team had deliberately recorded.

Salesforce permits the same contact to hold several roles on one opportunity, so all of them are preserved. Where someone is also marked as the primary contact, that flag is applied once per opportunity, as Salesforce requires.

Effect: 437 source rows produced 515 role records. Nothing was discarded.

Transformations applied

Data rarely moves across unchanged. These are the substantive conversions — the cases where a value in Salesforce deliberately differs from what's in Airtable.

Corrections made automatically

- **“Decomissioned” → “Decommissioned”** — a misspelling on 89 application records, corrected on load. Worth fixing at source so it isn't retyped.
- **“Gov?t Employees” → “Gov't Employees”** — 25 opportunity records have a question mark where an apostrophe should be. We confirmed this is genuinely stored that way in Airtable, not a glitch in our export.
- **Impediment categories** — “Issue on their end” and “Relationship Issue” were mapped to the equivalent Salesforce values.
- **Placeholder text removed** — “TBD” entries in web address fields were left blank rather than stored as if they were real links.

Values added to Salesforce so nothing was lost

- **Demographic Served** — Airtable uses 32 categories where Salesforce offered 6. We analysed each category by how recently it was used and added the 24 in active use within the last 18 months. Only 8 genuinely dormant categories were left out, affecting about 28 records rather than the ~400 a simple match would have dropped.
- **Contact roles** — “Day-to-Day POC” and “Senior POC” did not exist in Salesforce. Rather than forcing them into an approximate existing role, both were added.
- **Opportunity type and service level** — several values used throughout Airtable were not selectable on Login.gov opportunities. They were enabled, which also fixed revenue calculations that depend on them.

Fields deliberately not migrated

- **Calculated fields** — revenue totals, completion percentages and similar are recalculated by Salesforce from the underlying data, so writing them across would be overwritten anyway.
- **Market Segment** — set automatically from the related account.
- **Confirmed not needed** — Pilots, Usage Tracker Application Name, Vital Update %, and Issuer Strings, per earlier decisions.

Decisions taken during the load

Where the source data was ambiguous, we stopped rather than guessed. These are the calls that were made, and what each one cost or preserved.

The impediment named “None” was not migrated

One impediment record is named “None”, has no description, and is linked to 465 opportunities — more than five times any real impediment.

Loading it would have asserted that 297 opportunities *are* impeded, and because Salesforce totals blocked revenue per impediment, “None” would have appeared at the top of any “biggest blockers” report with a meaningless multi-million-dollar figure.

Effect: the 267 links from genuine impediments loaded and now show real figures. Reversing this decision takes one setting, not a rebuild.

Contacts without a name were loaded using their email address

Only 491 of 1,599 contacts have a name recorded. Salesforce requires a last name on every contact.

Rather than drop them or invent names, the email address is used as a visible placeholder, so the affected records are easy to find and correct later.

Effect: no contacts lost. About 971 look like real individuals whose name simply wasn't recorded; roughly 116 look like shared mailboxes.

Unmatched accounts were left for human review, not auto-created

169 Airtable account rows match no existing Salesforce account. Creating them automatically would have produced duplicate agency records, since spot-checking showed most are duplicate rows within Airtable rather than genuinely new organisations.

Effect: the largest remaining gap, but a reversible one — each resolved account releases its partner accounts, applications and links automatically.

What we need from you

In rough order of how many records each one unblocks.

HIGHEST IMPACT **Resolve the 169 unmatched accounts**
For each, tell us whether it duplicates an existing account (and which one is correct), is genuinely new, or is stale. This alone unblocks partner accounts, applications and several hundred links.

DECISION **Confirm what the “None” impediment means**
If it means “no impediment”, the cleanest fix is to remove those links in Airtable and delete the record.

DECISION **Service accounts and shared mailboxes**
Roughly 116 contacts are addresses like a help desk rather than people. Should they be contacts at all, or handled another way?

DATA ENTRY **Names for approximately 971 contacts**
The single highest-value cleanup available — ideally captured as first and last name rather than one free-text field.

DATA ENTRY **28 opportunities have no status**
and 16 have no account link, so they cannot be created in Salesforce.

DECISION **Confirm the record ownership rules**

Which account should own records whose Airtable owner has no Salesforce account; whether applications and contacts should take their account's owner; and whether existing accounts' ownership should be left alone. See "Record ownership" above.

PROVISIONING **Four opportunity leads have no Salesforce account**

Covering 285 opportunities between them. Two are already on this list from the partner account review, so resolving them fixes ownership in both places.

CLARIFY **Does "Technical Emails" mean "Technical POC"?**

716 contacts carry this subscription value, which has no Salesforce equivalent. We did not assume the two are the same.

CLARIFY **Seven opportunities are marked as both blocked by and requesting the same impediment.**

We recorded these as blocked, being the more severe reading.

Building the automated loading process

The goal is a repeatable process the team can re-run, not a one-time copy of the data. That distinction matters when reading the table above: a data set can be fully automated and still be mostly unloaded, because it is waiting on source data rather than on engineering. Applications is exactly that — built, tested and proven, but only 65% loaded.

Each data set moves through four stages. *Pull* reads Airtable directly through its API. *Transform* converts it and checks it against what Salesforce will actually accept. *Load* writes it. *Verified* means we confirmed the result in Salesforce rather than trusting the load report.

DATA SET	PULL	TRANSFORM	LOAD	VERIFIED	AUTOMATION NOTES
Accounts	✓	✓	✓	✓	Matches existing accounts rather than creating them, and writes anything it can't match confidently to a review file instead of guessing.
Partner Accounts	✓	✓	✓	✓	Re-runs pick up newly-resolved accounts automatically.
Contacts	✓	✓	✓	✓	Consolidates duplicate people across both source tables as part of the transform.
Opportunities	✓	✓	✓	✓	Required two Salesforce configuration changes, both applied.
Applications	✓	✓	✓	✓	Fully automated; the gap is source data, not the process.
Impediments	✓	✓	✓	✓	Complete apart from two blank source rows.

DATA SET	PULL	TRANSFORM	LOAD	VERIFIED	AUTOMATION NOTES
Application ↔ Contact links	✓	✓	✓	✓	Re-running can't create duplicates, by design.
Impediment ↔ Opportunity links	✓	✓	✓	✓	The “None” exclusion is a setting, not a code change.
Opportunity contact roles	✓	✓	✓	✓	Cannot use the usual matching key, so it compares against existing records instead. Re-running is still safe.
Meetings	✓	–	–	–	Data pulls fine. Deferred rather than blocked — the times don't exist in Airtable to load. See the known issue below.
Notes	✓	–	–	–	Deliberately last — a note must attach to a record that already exists.
Record ownership	✓	✓	–	–	Built once the four decisions were taken; applies at the next load.
Broker application links	✓	✓	–	–	Built. Generated automatically with the applications, loaded as a second step because Salesforce cannot link a record to another in the same batch.

WHAT "AUTOMATED" BUYS YOU

Every step re-reads Airtable and re-checks Salesforce each time it runs. That means the cleanup items on this page don't need an engineer to act on them individually — correcting the source data and re-running is enough to bring the newly-valid records across. Nothing has to be re-keyed by hand, and re-running cannot duplicate what is already loaded.

Record ownership

Raised by the Partnerships team and not yet built. Worth calling out separately because it affects who can see what, not just how reports read.

Every record the migration created is currently owned by the person who ran the load. Accounts are the exception — they were matched to records that already existed, so their original owners survived.

That matters more than it might appear: these objects use owner-based sharing, so the owner determines visibility, not just the “Owner” column on a report.

Proposed rule

Where Airtable records an owner and that person has a Salesforce account, assign the record to them. Where they don't, fall back to a single default owner.

This never leaves a record unowned and never invents an owner. Applied to today's data:

- **Opportunities** — 15 named leads. 11 have Salesforce accounts, covering **541 opportunities**. The other 4 cover 285, which would fall back.
- **Partner Accounts** — 5 of 7 owners resolve, covering 62 records.
- **Applications** — 5 of 7 resolve, covering 847 records.
- **Accounts and Contacts** — Airtable records no owner at all, so these need a decision rather than a rule.

Note: the four Opportunity leads without Salesforce accounts include two people already flagged on this list for the same reason — the same missing accounts block ownership in two places, so provisioning them fixes both.

Four questions before this can be built: which account should be the fallback owner; whether Applications should take their account's owner (Airtable has no application-specific owner); whether Contacts should do the same; and whether the migration should change ownership of accounts that already existed in Salesforce, or leave those untouched.

Known issue: meetings cannot be migrated as calendar events

This is the one part of the migration that cannot be solved by writing better code, and it needs a decision outside the migration itself.

Airtable records a meeting date, but no start time, end time or duration. Salesforce calendar events require them. There is no version of this transformation that doesn't invent information: a fixed start time would have 1,604 meetings all claiming to begin at 9am, and even an "all-day" event asserts something the source never recorded. Those times would then look exactly like real scheduling history to anyone reading the record later.

Decision: don't invent the times — recover the real ones

Rather than fabricate a schedule, the recommendation is to set up **Einstein Activity Capture**, the Salesforce feature that synchronises users' Google Calendar entries and associates them with Salesforce records. Once real calendar events are present, each Airtable meeting can be matched to the actual event it describes — using the date, who led it, who attended, and the meeting name — and the Airtable summary, agenda and action items attached to that real event.

Effect: the calendar supplies the times, Airtable supplies the meaning, and nothing is fabricated. **Meetings are deferred until this is in place** and are not part of the next load.

What this depends on, and what it will not recover

Two things should be understood before committing to this route, because they cap how much meeting history can realistically be recovered:

- **Roughly 3 in 10 meetings can never be matched, regardless of how well it works.** Calendar synchronisation works per person, from that person's own calendar. Only **1,315 of the 1,845** meetings were led by someone who has an active Salesforce user today; the remaining 530 were led by people who have left or were never given an account — one person alone led 182 of them. Those meetings have no calendar to synchronise from. *Giving those people Salesforce accounts, where they are still staff, raises this ceiling and also fixes the record-ownership gaps described above — the same short list of names appears in both.*
- **How far back the calendar history reaches is not yet known.** These meetings span February 2024 to August 2026 — 125 in 2024, 931 in 2025, 548 in 2026. Calendar synchronisation keeps a limited history and may only capture meetings from the day it is switched on. If it does not reach back, older meetings will never appear no matter how good the matching is.

> WHAT WE NEED FROM YOU
>
> A decision on whether to pursue Einstein Activity Capture, and who can confirm how far back its calendar history reaches in this organisation. Until that is answered we cannot say how many of the 1,845 meetings are recoverable — only that inventing meeting times is the wrong way to reach a higher number.
>
> We also need a decision on the meetings that will never match: discard them, or keep their summaries and action items as notes attached to the relevant agency or opportunity. The second option preserves the record of what was discussed without pretending a calendar entry existed.

Still to come

- **Record ownership** — *built, awaiting the next load.* All four decisions were taken: records are assigned to their Airtable owner where that person has an active Salesforce account, and otherwise to a named fallback owner rather than to whoever happens to run the load. Applications follow their partner account's owner; contacts follow their agency's owner.

- **Broker application relationships** — *built, awaiting the next load*. 63 of 70 links are ready; the rest depend on applications still blocked by the account duplicates described above.
- **Notes** — freeform Airtable columns with no dedicated Salesforce field (for example partner account descriptions and known blockers) become attached notes. This runs last, once every record it attaches to exists. It is the next piece of work.
- **Meetings** — deferred, see the known issue above.
- **A re-run after cleanup** — every load re-reads Airtable, so fixing the items above and re-running picks up the newly valid records with no additional engineering.

Counts were taken directly from the gsa-peo sandbox on 13 August 2026 and from the Airtable export pulled 12 August 2026. Because both systems remain in active use, figures will drift; this report is a snapshot, not a live view. Full field-by-field mapping detail and the complete list of data-quality items are maintained alongside the migration scripts.